

ch9exercises

Lauren Temple


3/2/2022

```
pacman::p_load(e1071, ROCR, ISLR2, caTools)
```

9.3 Here we explore the maximal margin classifier on a toy data set.


- a. We are given $n = 7$ observations in $p = 2$ dimensions. For each observation, there is an associated class label. Sketch the observations.

```
plot(-1:5, -1:5, type="n", xlab= 'X1', ylab= 'X2')
points(c(3, 2, 4, 1), c(4, 2, 4, 2), col= "red", pch= 19)
points(c(2, 4, 4), c(1, 3, 1), col= "blue", pch= 19)
```



- b. Sketch the optimal separating hyperplane, and provide the equation for this hyperplane (of the form (9.1)).

```
plot(-1:5, -1:5, type="n", xlab= 'X1', ylab= 'X2')
points(c(3, 2, 4, 1), c(4, 2, 4, 2), col= "red", pch= 19)
points(c(2, 4, 4), c(1, 3, 1), col= "blue", pch= 19)
abline(-0.5, 1) #y intercept=-0.5 and gradient=1.
```



$$y = mx + c$$

And given X_1 , X_2 , $m=1$ and $c=-0.5$, we have:

$$X_2 - X_1 + 0.5 = 0$$


- c. Describe the classification rule for the maximal margin classifier. It should be something along the lines of “Classify to Red if $0 + 1X_1 + 2X_2 > 0$, and classify to Blue otherwise.” Provide the values for 0, 1, and 2.

Class Red if $X_2 - X_1 + 0.5 > 0$ and Blue otherwise. $\beta_0 = 0.5; \beta_1 = -1; \beta_2 = 1$

- d. On your sketch, indicate the margin for the maximal margin hyperplane.


```
plot(-1:5, -1:5, type="n", xlab= 'X1', ylab= 'X2')
points(c(3, 2, 4, 1), c(4, 2, 4, 2), col= "red", pch= 19)
points(c(2, 4, 4), c(1, 3, 1), col= "blue", pch= 19)
abline(-0.5, 1) #y intercept=-0.5 and gradient=1.
```

```
abline(-1, 1, lty= 'dotted')
abline(0, 1, lty= 'dotted')
```



e. Indicate the support vectors for the maximal margin classifier.

```
knitr::include_graphics("9.3.jpg")
```




- f. Argue that a slight movement of the seventh observation would not affect the maximal margin hyperplane.

A slight movement of the seventh observation would not have an effect on the maximal margin hyperplane since it is located outside of the margin so a slight movement.

- g. Sketch a hyperplane that is not the optimal separating hyperplane, and provide the equation for this hyperplane.


```
plot(-1:5, -1:5, type="n", xlab= 'X1', ylab= 'X2')
points(c(3, 2, 4, 1), c(4, 2, 4, 2), col= "red", pch= 19)
points(c(2, 4, 4), c(1, 3, 1), col= "blue", pch= 19)
abline(-0.5, 1) #y intercept=-0.5 and gradient=1.
abline(0,0.8, col="green") #not optimal
abline(-1, 1, lty= 'dotted')
abline(0, 1, lty= 'dotted')
```



- h. Draw an additional observation on the plot so that the two classes are no longer separable by a hyperplane.

```
plot(-1:5, -1:5, type="n", xlab= 'X1', ylab= 'X2')
points(c(3, 2, 4, 1), c(4, 2, 4, 2), col= "red", pch= 19)
points(c(2, 4, 4), c(1, 3, 1), col= "blue", pch= 19)
points(2, 3, col= "blue", pch= 17)
```

```
abline(-0.5, 1) #y intercept=-0.5 and gradient=1.
abline(-1, 1, lty= 'dotted')
abline(0, 1, lty= 'dotted')
```



9.5 We have seen that we can fit an SVM with a non-linear kernel in order to perform classification using a non-linear decision boundary. We will now see that we can also obtain a non-linear decision boundary by performing logistic regression using non-linear transformations of the features.


- Generate a data set with $n = 500$ and $p = 2$, such that the observations belong to two classes with a quadratic decision boundary between them.

```
set.seed(42)
x1 <- runif(500)-0.5
x2 <- runif(500)-0.5
y <- 1*(x1^2-x2^2 > 0)
df <- data.frame(x1= x1, x2= x2, y=as.factor(y))
head(df)
```

```
##           x1           x2 y
## 1  0.41480604 -0.36349479 1
## 2  0.43707541 -0.32286359 1
## 3 -0.21386047  0.01956045 1
## 4  0.33044763  0.31112079 1
## 5  0.14174552 -0.38463799 0
## 6  0.01909595  0.39342179 0
```

- b. Plot the observations, colored according to their class labels. Your plot should display X1 on the x-axis, and X2 on the yaxis.

```
plot(x1, x2, col= (2 - y))
```



- c. Fit a logistic regression model to the data, using X1 and X2 as predictors.

```
glm.fit <- glm(y~ x1+x2, data= df, family= 'binomial')
```


- d. Apply this model to the training data in order to obtain a predicted class label for each training observation. Plot the observations, colored according to the predicted class labels. The decision boundary should be linear.

```
glm.probs <- predict(glm.fit, newdata= df, type= 'response')
glm.preds <- rep(0, 500)
glm.preds[glm.probs>0.5] = 1

table(preds= glm.preds, truth= df$y)
```

```
##      truth
## preds   0   1
##      0  24  22
##      1 213 241
```

```
plot(x1, x2, col= 2-glm.preds)
```



- e. Now fit a logistic regression model to the data using non-linear functions of X_1 and X_2 as predictors (e.g. X_2 , $X_1 \times X_2$, $\log(X_2)$, and so forth).

```
glm.fit <- glm(y~I(x1^2)+I(x2^2), data = df, family = 'binomial')
```

```
## Warning: glm.fit: algorithm did not converge
```

```
## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred
```


- f. Apply this model to the training data in order to obtain a predicted class label for each training observation. Plot the observations, colored according to the predicted class labels. The decision boundary should be obviously non-linear. If it is not, then repeat (a)-(e) until you come up with an example in which the predicted class labels are obviously non-linear.

```
glm.probs <- predict(glm.fit, newdata= df, type= 'response')
glm.preds <- rep(0,500)
glm.preds[glm.probs>0.5] = 1
table(preds= glm.preds, truth=df$y)
```



```
##      truth
## preds  0   1
##      0 237   0
##      1   0 263
```

```
plot(x1, x2, col=2-glm.preds)
```




- g. Fit a support vector classifier to the data with X1 and X2 as predictors. Obtain a class prediction for each training observation. Plot the observations, colored according to the predicted class labels.

```
tune.out <- tune(svm, y~., data=df, kernel='linear', ranges= list(cost=c(0.001, 0.01, 0.1, 1, 5, 10, 100),
bestmod <- tune.out$best.model
```

```
ypred <- predict(bestmod, newdata= df, type='response')
table(predict=ypred, truth=df$y)
```

```
##      truth
## predict  0   1
##      0   0   0
##      1 237 263
```

```
plot(x1, x2, col=ypred)
```




- h. Fit a SVM using a non-linear kernel to the data. Obtain a class prediction for each training observation. Plot the observations, colored according to the predicted class labels.

```
tune.out <- tune(svm, y~., data=df, kernel='radial', ranges= list(cost=c(0.1, 1, 10, 100, 1000), gamma=
bestmod <- tune.out$best.model
```

```
ypred <- predict(bestmod, newdata=df, type='response')
table(predict= ypred, truth= df$y)
```

```
##      truth
## predict  0   1
##      0 236   0
##      1   1 263
```

```
plot(x1, x2, col= ypred)
```



i. Comment on your results.

Logistic regression with non-linear terms and SVM with a radial kernel outperform the models using linear terms. The Logistic regression with non-linear terms and SVM with a radial kernel models achieve almost perfect accuracy in predicting the class of the training observations.

9.7 In this problem, you will use support vector approaches in order to predict whether a given car gets high or low gas mileage based on the Auto data set.

a. Create a binary variable that takes on a 1 for cars with gas mileage above the median, and a 0 for cars with gas mileage below the median.

```
set.seed(42)
auto.length <- length(Auto$mpg)
mpg.median <- median(Auto$mpg)
mpg01 <- rep(NA, auto.length)

for(i in 1:auto.length) if (Auto$mpg[i] > mpg.median) mpg01[i]= 1 else mpg01[i]= 0

auto.df <- Auto
auto.df$mpg01 <- as.factor(mpg01)
```

b. Fit a support vector classifier to the data with various values of cost, in order to predict whether a car gets high or low gas mileage. Report the cross-validation errors associated with different values of this parameter. Comment on your results. Note you will need to fit the classifier without the gas mileage variable to produce sensible results.

```
linear.tune <- tune(svm, mpg01~., data= auto.df, kernel= 'linear', ranges= list(cost=c(0.001, 0.01, 0.1),
summary(linear.tune)
```

```
##
## Parameter tuning of 'svm':
##
## - sampling method: 10-fold cross validation
##
## - best parameters:
##   cost
##     1
##
## - best performance: 0.0075
##
## - Detailed performance results:
##   cost      error dispersion
## 1 1e-03 0.09955128 0.05722671
## 2 1e-02 0.07391026 0.05301906
## 3 1e-01 0.04326923 0.03594224
## 4 1e+00 0.00750000 0.01687371
## 5 5e+00 0.01519231 0.01760469
## 6 1e+01 0.01775641 0.01700310
## 7 1e+02 0.03044872 0.02842322
## 8 1e+03 0.03044872 0.02842322
```

Training CV error decreases as cost increases with minimum at cost=1 then it begins to increase.

- c. Now repeat (b), this time using SVMs with radial and polynomial basis kernels, with different values of gamma and degree and cost. Comment on your results.

```
set.seed(42)
radial.tune <- tune(svm, mpg01~., data= auto.df, kernel='radial', ranges= list(cost=c(0.1, 1, 10, 100, 1000),
radial.tune$best.parameters
```

```
##   cost gamma
## 3    10    0.5
```

```
radial.tune$best.performance
```

```
## [1] 0.05070513
```

Training CV is lowest for a radial model with cost=10 and gamma=0.5

```
set.seed(42)
poly.tune <- tune(svm, mpg01~., data= auto.df, kernel='polynomial', ranges=list(cost=c(0.1, 1, 10, 100, 1000),
poly.tune$best.parameters
```

```
##   cost degree
## 5 1000      1
```

```
poly.tune$best.performance
```

```
## [1] 0.007564103
```

The training error is minimized with $\text{cost} = 1000$ and $\text{degree} = 1$

- d. Make some plots to back up your assertions in (b) and (c).

9.8 This problem involves the OJ data set which is part of the ISLR2 package.

```
data(OJ)
OJ <- OJ
```

- (a) Create a training set containing a random sample of 800 observations, and a test set containing the remaining observations.

```
sample.data <- sample.split(OJ$Purchase, SplitRatio = 800/length(OJ$Purchase))
train.set <- subset(OJ, sample.data==T)
test.set <- subset(OJ, sample.data==F)
```

- b. Fit a support vector classifier to the training data using $\text{cost} = 0.01$, with Purchase as the response and the other variables as predictors. Use the `summary()` function to produce summary statistics, and describe the results obtained.

```
svm.fit <- svm(Purchase~., data= train.set, kernel='linear', cost=0.01)
summary(svm.fit)
```

```
##
## Call:
## svm(formula = Purchase ~ ., data = train.set, kernel = "linear",
##      cost = 0.01)
##
##
## Parameters:
##   SVM-Type:  C-classification
##   SVM-Kernel: linear
##      cost:  0.01
##
## Number of Support Vectors:  431
##
## ( 216 215 )
##
##
## Number of Classes:  2
##
## Levels:
##  CH MM
```

There are a large number of support vectors out of 800 observations, 216 belong to CH, 215 belong to MM.

- c. What are the training and test error rates?

```
svm.pred <- predict(svm.fit, train.set)
er1 <- table(predict= svm.pred, truth= train.set$Purchase)
1 - sum(diag(er1)) / sum(er1)
```

```
## [1] 0.16625
```

```
svm.pred <- predict(svm.fit, test.set)
er2 <- table(predict= svm.pred, truth= test.set$Purchase)
1 - sum(diag(er2)) / sum(er2)
```

```
## [1] 0.1777778
```

d. Use the `tune()` function to select an optimal cost. Consider values in the range 0.01 to 10.

```
set.seed(42)
tune.out <- tune(svm, Purchase~., data=train.set, kernel= 'linear', range= list(cost=c(0.01, 0.1, 0.5, 1, 10)))
```

e. Compute the training and test error rates using this new value for cost.

```
svm.pred <- predict(tune.out$best.model, train.set)
er <- table(predict= svm.pred, truth= train.set$Purchase)

1 - sum(diag(er)) / sum(er)
```

```
## [1] 0.1625
```

```
svm.pred <- predict(tune.out$best.model, test.set)
er <- table(predict= svm.pred, truth= test.set$Purchase)

1 - sum(diag(er)) / sum(er)
```

```
## [1] 0.1814815
```

Using the optimal value of cost results in a slightly increased test error rate.

f. Repeat parts (b) through (e) using a support vector machine with a radial kernel. Use the default value for gamma.

```
set.seed(42)
tune.out <- tune(svm, Purchase~., data= train.set, kernel='radial', ranges=list(cost=c(0.01, 0.1, 0.5, 1, 10)))
```

```
svm.pred <- predict(tune.out$best.model, train.set)
er <- table(predict= svm.pred, truth= train.set$Purchase)

1 - sum(diag(er)) / sum(er)
```

```
## [1] 0.14
```

```
svm.pred <- predict(tune.out$best.model, test.set)
er <- table(predict= svm.pred, truth= test.set$Purchase)

1 - sum(diag(er)) / sum(er)
```

```
## [1] 0.2111111
```

- g. Repeat parts (b) through (e) using a support vector machine with a polynomial kernel. Set degree = 2.

```
set.seed(42)
tune.out <- tune(svm, Purchase~., data= train.set, kernel= 'polynomial', ranges= list(cost=c(0.01, 0.1,
```

```
svm.pred <- predict(tune.out$best.model, train.set)
er <- table(predict= svm.pred, truth= train.set$Purchase)

1 - sum(diag(er)) / sum(er)
```

```
## [1] 0.14125
```

```
svm.pred <- predict(tune.out$best.model, test.set)
er <- table(predict= svm.pred, truth= test.set$Purchase)

1 - sum(diag(er)) / sum(er)
```

```
## [1] 0.2111111
```

- h. Overall, which approach seems to give the best results on this data?

Overall the radial kernel seems to be producing minimum misclassification error on both training and test data.